# C2_W4_Assignment

October 31, 2020

# 1 Assignment 4: Word Embeddings

Welcome to the fourth (and last) programming assignment of Course 2!

In this assignment, you will practice how to compute word embeddings and use them for sentiment analysis. - To implement sentiment analysis, you can go beyond counting the number of positive words and negative words. - You can find a way to represent each word numerically, by a vector. - The vector could then represent syntactic (i.e. parts of speech) and semantic (i.e. meaning) structures.

In this assignment, you will explore a classic way of generating word embeddings or representations. - You will implement a famous model called the continuous bag of words (CBOW) model.

By completing this assignment you will:

- Train word vectors from scratch.
- Learn how to create batches of data.
- Understand how backpropagation works.
- Plot and visualize your learned word vectors.

Knowing how to train these models will give you a better understanding of word vectors, which are building blocks to many applications in natural language processing.

## 1.1 Outline

- Section **??**
- Section **??**

    - Section **??**
        * Section **??**
    - Section **??**
        * Section **??**
    - Section **??**
        * Section **??**
    - Section **??**
    - Section **??**
        * Section **??**
    - Section **??**

* Section **??**

- Section **??**

# 1. The Continuous bag of words model
Let's take a look at the following sentence: >**'I am happy because I am learning'**.

- In continuous bag of words (CBOW) modeling, we try to predict the center word given a few context words (the words around the center word).
- For example, if you were to choose a context half-size of say $C = 2$, then you would try to predict the word **happy** given the context that includes 2 words before and 2 words after the center word:

$C$ words before: [I, am]

$C$ words after: [because, I]

- In other words:

$$context = [I, am, because, I]$$

$$target = happy$$

The structure of your model will look like this:
Figure 1
Where $\bar{x}$ is the average of all the one hot vectors of the context words.
Figure 2
Once you have encoded all the context words, you can use $\bar{x}$ as the input to your model.
The architecture you will be implementing is as follows:

$$h = W_1\, X + b_1 \tag{1}$$
$$a = ReLU(h) \tag{2}$$
$$z = W_2\, a + b_2 \tag{3}$$
$$\hat{y} = softmax(z) \tag{4}$$
$$\tag{1}$$

```
In [ ]: # Import Python libraries and helper functions (in utils2)
        import nltk
        from nltk.tokenize import word_tokenize
        import numpy as np
        from collections import Counter
        from utils2 import sigmoid, get_batches, compute_pca, get_dict
```

```
In [ ]: # Download sentence tokenizer
        nltk.data.path.append('.')
```

2

```
In [ ]: # Load, tokenize and process the data
        import re                                                     #  Load the Regex-
        with open('shakespeare.txt') as f:
            data = f.read()                                           #  Read in the dat
        data = re.sub(r'[,!?;-]', '.',data)                           #  Punktuations ar
        data = nltk.word_tokenize(data)                               #  Tokenize string
        data = [ ch.lower() for ch in data if ch.isalpha() or ch == '.']   #  Lower case and
        print("Number of tokens:", len(data),'\n', data[:15])         #  print data samp

In [ ]: # Compute the frequency distribution of the words in the dataset (vocabulary)
        fdist = nltk.FreqDist(word for word in data)
        print("Size of vocabulary: ",len(fdist) )
        print("Most frequent tokens: ",fdist.most_common(20) ) # print the 20 most frequent wo
```

**Mapping words to indices and indices to words**   We provide a helper function to create a dictionary that maps words to indices and indices to words.

```
In [ ]: # get_dict creates two dictionaries, converting words to indices and viceversa.
        word2Ind, Ind2word = get_dict(data)
        V = len(word2Ind)
        print("Size of vocabulary: ", V)

In [ ]: # example of word to index mapping
        print("Index of the word 'king' :  ",word2Ind['king'] )
        print("Word which has index 2743:  ",Ind2word[2743] )
```

# 2 Training the Model

### 1.1.1   Initializing the model

You will now initialize two matrices and two vectors. - The first matrix ($W_1$) is of dimension $N \times V$, where $V$ is the number of words in your vocabulary and $N$ is the dimension of your word vector. - The second matrix ($W_2$) is of dimension $V \times N$. - Vector $b_1$ has dimensions $N \times 1$ - Vector $b_2$ has dimensions $V \times 1$. - $b_1$ and $b_2$ are the bias vectors of the linear layers from matrices $W_1$ and $W_2$.

The overall structure of the model will look as in Figure 1, but at this stage we are just initializing the parameters.

### Exercise 01 Please use numpy.random.rand to generate matrices that are initialized with random values from a uniform distribution, ranging between 0 and 1.

**Note:** In the next cell you will encounter a random seed. Please **DO NOT** modify this seed so your solution can be tested correctly.

```
In [ ]: # UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: initialize_model
        def initialize_model(N,V, random_seed=1):
            '''
            Inputs:
                N:   dimension of hidden vector
                V:   dimension of vocabulary
                random_seed: random seed for consistent results in the unit tests
```

3

```
    Outputs:
        W1, W2, b1, b2: initialized weights and biases
    '''

    np.random.seed(random_seed)

    ### START CODE HERE (Replace instances of 'None' with your code) ###
    # W1 has shape (N,V)
    W1 = np.random.rand(N,V)
    # W2 has shape (V,N)
    W2 = np.random.rand(V,N)
    # b1 has shape (N,1)
    b1 = np.random.rand(N,1)
    # b2 has shape (V,1)
    b2 = np.random.rand(V,1)
    ### END CODE HERE ###

    return W1, W2, b1, b2
```

```
In [ ]: # Test your function example.
        tmp_N = 4
        tmp_V = 10
        tmp_W1, tmp_W2, tmp_b1, tmp_b2 = initialize_model(tmp_N,tmp_V)
        assert tmp_W1.shape == ((tmp_N,tmp_V))
        assert tmp_W2.shape == ((tmp_V,tmp_N))
        print(f"tmp_W1.shape: {tmp_W1.shape}")
        print(f"tmp_W2.shape: {tmp_W2.shape}")
        print(f"tmp_b1.shape: {tmp_b1.shape}")
        print(f"tmp_b2.shape: {tmp_b2.shape}")
```

**Expected Output**

```
tmp_W1.shape: (4, 10)
tmp_W2.shape: (10, 4)
tmp_b1.shape: (4, 1)
tmp_b2.shape: (10, 1)
```

### 2.1 Softmax Before we can start training the model, we need to implement the softmax function as defined in equation 5:

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{i=0}^{V-1} e^{z_i}} \tag{5}$$

- Array indexing in code starts at 0.
- $V$ is the number of words in the vocabulary (which is also the number of rows of $z$).
- $i$ goes from 0 to $|V|$ - 1.

### Exercise 02 **Instructions**: Implement the softmax function below.

- Assume that the input $z$ to softmax is a 2D array

4

- Each training example is represented by a column of shape (V, 1) in this 2D array.
- There may be more than one column, in the 2D array, because you can put in a batch of examples to increase efficiency. Let's call the batch size lowercase $m$, so the $z$ array has shape (V, m)
- When taking the sum from $i = 1 \cdots V - 1$, take the sum for each column (each example) separately.

Please use - numpy.exp - numpy.sum (set the axis so that you take the sum of each column in z)

```
In [ ]: # UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: softmax
        def softmax(z):
            '''
            Inputs:
                z: output scores from the hidden layer
            Outputs:
                yhat: prediction (estimate of y)
            '''

            ### START CODE HERE (Replace instances of 'None' with your own code) ###

            # Calculate yhat (softmax)
            e_z = np.exp(z)
            yhat = e_z/np.sum(e_z,axis=0)

            ### END CODE HERE ###

            return yhat
```

```
In [ ]: # Test the function
        tmp = np.array([[1,2,3],
                        [1,1,1]
                       ])
        tmp_sm = softmax(tmp)
        display(tmp_sm)
```

**Expected Ouput**

```
array([[0.5       , 0.73105858, 0.88079708],
       [0.5       , 0.26894142, 0.11920292]])
```

### 2.2 Forward propagation
### Exercise 03 Implement the forward propagation $z$ according to equations (1) to (3).

$$h = W_1 \, X + b_1 \tag{1}$$
$$a = ReLU(h) \tag{2}$$
$$z = W_2 \, a + b_2 \tag{3}$$
$$\tag{2}$$

5

For that, you will use as activation the Rectified Linear Unit (ReLU) given by:

$$f(h) = \max(0, h) \tag{6}$$

Hints

You can use numpy.maximum(x1,x2) to get the maximum of two values

Use numpy.dot(A,B) to matrix multiply A and B

```python
In [ ]: # UNQ_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: forward_prop
        def forward_prop(x, W1, W2, b1, b2):
            '''
            Inputs:
                x:  average one hot vector for the context
                W1, W2, b1, b2:  matrices and biases to be learned
             Outputs:
                z:  output score vector
            '''

            ### START CODE HERE (Replace instances of 'None' with your own code) ###

            # Calculate h
            h = np.dot(W1,x)+b1

            # Apply the relu on h (store result in h)
            h = np.maximum(0,h)

            # Calculate z
            z = np.dot(W2,h)+b2

            ### END CODE HERE ###

            return z, h
```

```python
In [ ]: # Test the function

        # Create some inputs
        tmp_N = 2
        tmp_V = 3
        tmp_x = np.array([[0,1,0]]).T
        tmp_W1, tmp_W2, tmp_b1, tmp_b2 = initialize_model(N=tmp_N,V=tmp_V, random_seed=1)

        print(f"x has shape {tmp_x.shape}")
        print(f"N is {tmp_N} and vocabulary size V is {tmp_V}")

        # call function
        tmp_z, tmp_h = forward_prop(tmp_x, tmp_W1, tmp_W2, tmp_b1, tmp_b2)
        print("call forward_prop")
        print()
```

```python
        # Look at output
        print(f"z has shape {tmp_z.shape}")
        print("z has values:")
        print(tmp_z)

        print()

        print(f"h has shape {tmp_h.shape}")
        print("h has values:")
        print(tmp_h)
```

**Expected output**

```
x has shape (3, 1)
N is 2 and vocabulary size V is 3
call forward_prop

z has shape (3, 1)
z has values:
[[0.55379268]
 [1.58960774]
 [1.50722933]]

h has shape (2, 1)
h has values:
[[0.92477674]
 [1.02487333]]
```

## 2.3 Cost function

- We have implemented the *cross-entropy* cost function for you.

```python
In [ ]: # compute_cost: cross-entropy cost functioN
        def compute_cost(y, yhat, batch_size):
            # cost function
            logprobs = np.multiply(np.log(yhat),y) + np.multiply(np.log(1 - yhat), 1 - y)
            cost = - 1/batch_size * np.sum(logprobs)
            cost = np.squeeze(cost)
            return cost
```

```python
In [ ]: # Test the function
        tmp_C = 2
        tmp_N = 50
        tmp_batch_size = 4
        tmp_word2Ind, tmp_Ind2word = get_dict(data)
        tmp_V = len(word2Ind)

        tmp_x, tmp_y = next(get_batches(data, tmp_word2Ind, tmp_V,tmp_C, tmp_batch_size))
```

```python
        print(f"tmp_x.shape {tmp_x.shape}")
        print(f"tmp_y.shape {tmp_y.shape}")

        tmp_W1, tmp_W2, tmp_b1, tmp_b2 = initialize_model(tmp_N,tmp_V)

        print(f"tmp_W1.shape {tmp_W1.shape}")
        print(f"tmp_W2.shape {tmp_W2.shape}")
        print(f"tmp_b1.shape {tmp_b1.shape}")
        print(f"tmp_b2.shape {tmp_b2.shape}")

        tmp_z, tmp_h = forward_prop(tmp_x, tmp_W1, tmp_W2, tmp_b1, tmp_b2)
        print(f"tmp_z.shape: {tmp_z.shape}")
        print(f"tmp_h.shape: {tmp_h.shape}")

        tmp_yhat = softmax(tmp_z)
        print(f"tmp_yhat.shape: {tmp_yhat.shape}")

        tmp_cost = compute_cost(tmp_y, tmp_yhat, tmp_batch_size)
        print("call compute_cost")
        print(f"tmp_cost {tmp_cost:.4f}")
```

**Expected output**

```
tmp_x.shape (5778, 4)
tmp_y.shape (5778, 4)
tmp_W1.shape (50, 5778)
tmp_W2.shape (5778, 50)
tmp_b1.shape (50, 1)
tmp_b2.shape (5778, 1)
tmp_z.shape: (5778, 4)
tmp_h.shape: (50, 4)
tmp_yhat.shape: (5778, 4)
call compute_cost
tmp_cost 9.9560
```

## 2.4 Training the Model - Backpropagation
### Exercise 04 Now that you have understood how the CBOW model works, you will train it. You created a function for the forward propagation. Now you will implement a function that computes the gradients to backpropagate the errors.

```python
In [ ]: # UNQ_C4 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: back_prop
        def back_prop(x, yhat, y, h, W1, W2, b1, b2, batch_size):
            '''
            Inputs:
                x:  average one hot vector for the context
                yhat: prediction (estimate of y)
                y:  target vector
                h:  hidden vector (see eq. 1)
```

8

```
        W1, W2, b1, b2:  matrices and biases
        batch_size: batch size
     Outputs:
        grad_W1, grad_W2, grad_b1, grad_b2:  gradients of matrices and biases
     '''
    ### START CODE HERE (Replace instanes of 'None' with your code) ###

    # Compute l1 as W2^T (Yhat - Y)
    # Re-use it whenever you see W2^T (Yhat - Y) used to compute a gradient
    l1 = np.dot(W2.T,(yhat-y))
    # Apply relu to l1
    l1 = np.max(0,l1)
    # Compute the gradient of W1
    grad_W1 = (1/batch_size)*np.dot(l1,x.T)
    # Compute the gradient of W2
    grad_W2 = (1/batch_size)*np.dot(yhat-y,h.T)
    # Compute the gradient of b1
    grad_b1 = np.sum((1/batch_size)*np.dot(l1,x.T),axis=1,keepdims=True)
    # Compute the gradient of b2
    grad_b2 = np.sum((1/batch_size)*np.dot(yhat-y,h.T),axis=1,keepdims=True)
    ### END CODE HERE ###

    return grad_W1, grad_W2, grad_b1, grad_b2
```

In [ ]:
```
# Test the function
tmp_C = 2
tmp_N = 50
tmp_batch_size = 4
tmp_word2Ind, tmp_Ind2word = get_dict(data)
tmp_V = len(word2Ind)

# get a batch of data
tmp_x, tmp_y = next(get_batches(data, tmp_word2Ind, tmp_V,tmp_C, tmp_batch_size))

print("get a batch of data")
print(f"tmp_x.shape {tmp_x.shape}")
print(f"tmp_y.shape {tmp_y.shape}")

print()
print("Initialize weights and biases")
tmp_W1, tmp_W2, tmp_b1, tmp_b2 = initialize_model(tmp_N,tmp_V)

print(f"tmp_W1.shape {tmp_W1.shape}")
print(f"tmp_W2.shape {tmp_W2.shape}")
print(f"tmp_b1.shape {tmp_b1.shape}")
print(f"tmp_b2.shape {tmp_b2.shape}")

print()
```

```python
        print("Forwad prop to get z and h")
        tmp_z, tmp_h = forward_prop(tmp_x, tmp_W1, tmp_W2, tmp_b1, tmp_b2)
        print(f"tmp_z.shape: {tmp_z.shape}")
        print(f"tmp_h.shape: {tmp_h.shape}")

        print()
        print("Get yhat by calling softmax")
        tmp_yhat = softmax(tmp_z)
        print(f"tmp_yhat.shape: {tmp_yhat.shape}")

        tmp_m = (2*tmp_C)
        tmp_grad_W1, tmp_grad_W2, tmp_grad_b1, tmp_grad_b2 = back_prop(tmp_x, tmp_yhat, tmp_y,

        print()
        print("call back_prop")
        print(f"tmp_grad_W1.shape {tmp_grad_W1.shape}")
        print(f"tmp_grad_W2.shape {tmp_grad_W2.shape}")
        print(f"tmp_grad_b1.shape {tmp_grad_b1.shape}")
        print(f"tmp_grad_b2.shape {tmp_grad_b2.shape}")
```

**Expected output**

```
get a batch of data
tmp_x.shape (5778, 4)
tmp_y.shape (5778, 4)

Initialize weights and biases
tmp_W1.shape (50, 5778)
tmp_W2.shape (5778, 50)
tmp_b1.shape (50, 1)
tmp_b2.shape (5778, 1)

Forwad prop to get z and h
tmp_z.shape: (5778, 4)
tmp_h.shape: (50, 4)

Get yhat by calling softmax
tmp_yhat.shape: (5778, 4)

call back_prop
tmp_grad_W1.shape (50, 5778)
tmp_grad_W2.shape (5778, 50)
tmp_grad_b1.shape (50, 1)
tmp_grad_b2.shape (5778, 1)
```

## Gradient Descent
### Exercise 05 Now that you have implemented a function to compute the gradients, you will implement batch gradient descent over your training set.

**Hint:** For that, you will use `initialize_model` and the `back_prop` functions which you just created (and the `compute_cost` function). You can also use the provided `get_batches` helper function:

```
for x, y in get_batches(data, word2Ind, V, C, batch_size):
...
```

Also: print the cost after each batch is processed (use batch size = 128)

```python
In [ ]: # UNQ_C5 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: gradient_descent
        def gradient_descent(data, word2Ind, N, V, num_iters, alpha=0.03):

            '''
            This is the gradient_descent function

              Inputs:
                data:      text
                word2Ind:  words to Indices
                N:         dimension of hidden vector
                V:         dimension of vocabulary
                num_iters: number of iterations
              Outputs:
                W1, W2, b1, b2:  updated matrices and biases

            '''
            W1, W2, b1, b2 = initialize_model(N,V, random_seed=282)
            batch_size = 128
            iters = 0
            C = 2
            for x, y in get_batches(data, word2Ind, V, C, batch_size):
                ### START CODE HERE (Replace instances of 'None' with your own code) ###
                # Get z and h
                z, h = forward_prop(x, W1, W2, b1, b2)
                # Get yhat
                yhat = softmax(z)
                # Get cost
                cost = compute_cost(y, yhat, batch_size)
                if ( (iters+1) % 10 == 0):
                    print(f"iters: {iters + 1} cost: {cost:.6f}")
                # Get gradients
                grad_W1, grad_W2, grad_b1, grad_b2 = back_prop(x, yhat, y, h, W1, W2, b1, b2, l

                # Update weights and biases
                W1 -= alpha*grad_W1
                W2 -= alpha*grad_W2
                b1 -= alpha*grad_b1
                b2 -= alpha*grad_b2

                ### END CODE HERE ###
```

```
            iters += 1
            if iters == num_iters:
                break
            if iters % 100 == 0:
                alpha *= 0.66


        return W1, W2, b1, b2
```

In [ ]: `# test your function`
```
        C = 2
        N = 50
        word2Ind, Ind2word = get_dict(data)
        V = len(word2Ind)
        num_iters = 150
        print("Call gradient_descent")
        W1, W2, b1, b2 = gradient_descent(data, word2Ind, N, V, num_iters)
```

**Expected Output**

```
iters: 10 cost: 0.789141
iters: 20 cost: 0.105543
iters: 30 cost: 0.056008
iters: 40 cost: 0.038101
iters: 50 cost: 0.028868
iters: 60 cost: 0.023237
iters: 70 cost: 0.019444
iters: 80 cost: 0.016716
iters: 90 cost: 0.014660
iters: 100 cost: 0.013054
iters: 110 cost: 0.012133
iters: 120 cost: 0.011370
iters: 130 cost: 0.010698
iters: 140 cost: 0.010100
iters: 150 cost: 0.009566
```

Your numbers may differ a bit depending on which version of Python you're using.
## 3.0 Visualizing the word vectors
In this part you will visualize the word vectors trained using the function you just coded above.

In [ ]: `# visualizing the word vectors here`
```
        from matplotlib import pyplot
        %config InlineBackend.figure_format = 'svg'
        words = ['king', 'queen','lord','man', 'woman','dog','wolf',
                'rich','happy','sad']

        embs = (W1.T + W2)/2.0

        # given a list of words and the embeddings, it returns a matrix with all the embedding
        idx = [word2Ind[word] for word in words]
```

```
        X = embs[idx, :]
        print(X.shape, idx)   # X.shape:  Number of words of dimension N each

In [ ]: result= compute_pca(X, 2)
        pyplot.scatter(result[:, 0], result[:, 1])
        for i, word in enumerate(words):
            pyplot.annotate(word, xy=(result[i, 0], result[i, 1]))
        pyplot.show()
```

You can see that man and king are next to each other. However, we have to be careful with the interpretation of this projected word vectors, since the PCA depends on the projection – as shown in the following illustration.

```
In [ ]: result= compute_pca(X, 4)
        pyplot.scatter(result[:, 3], result[:, 1])
        for i, word in enumerate(words):
            pyplot.annotate(word, xy=(result[i, 3], result[i, 1]))
        pyplot.show()
```